J. CHRISTOPHER MALONEY

SAVING PSYCHOLOGICAL SOLIPSISM*

(Received in revised form 10 June, 1989)

Bert, as you may know, prefers the familiar: saddle shoes, paper clips and the classic computationalist edition of the Representational Theory of the Mind bound in Methodological Solipsism [Fodor, 1987, 1982, 1980, 1975; Maloney, 1989]. Recently, Bert has been distracted and distressed. Oh, he still sports his saddle shoes, and his paper clip collection continues to expand. It's just that Lynne Rudder Baker [1987a, 1987b, 1986, 1985], following the lead of Tyler Burge [1986, 1982a, 1982b, 1979] and perhaps Stephen Stich [1983], has been badgering him about that other stuff. Let's try to get Bert's blood pressure back to normal.

It all started when Bert, physicalist that he is, began to wonder how it was that psychological states — beliefs and desires paradigmatically — could be intentional states whose contents figure so centrally in causal explanations of intelligent behavior. Bert takes it for granted that Ernie, for example, generally acts as he does because he believes and desires as he does. It has long been Bert's practice, when attempting to say just which of Ernie's psychological states cause Ernie's behavior, to identify Ernie's efficacious psychological states by specifying their contents. Psychological states, Bert supposes, are relations to mental representations. These representations have content, which they bestow on those psychological states in which they occur as relata. So, since the behavior that results from psychological states typically accords with the contents of those states, Bert reasons that mental representations must carry the causal charge of psychological states and, therefore be physical structures.

Evidently, then, the physical form and content of a representation would need to covary. Nonetheless, Bert recognizes that psychological types are multiply realizable and, thus, not covariant with, say, neurophysical types. The correspondence between the physical and the

Philosophical Studies **61**: 267—283, 1991.

psychological must be like the variation between syntax and semantics in a language, supposing that mental representations collectively function as elements in a mental language — Mentalese. Consequently, agents of quite different physical kinds could realize, in different kinds of internal physical states, mental states of the same syntactic type. Thus, Bert hypothesized, psychological content must ride on the physical just as meaning would supervene on syntactic structure in a thoroughly *unambiguous* public language. Occurrences, then, of a specified syntactic type in Mentalese would thereby be occurrences of a given psychological state with some determinate content. And occurrences of that psychological state would cause the actions that they do by the grace of having the syntactical properties that they do. That would seem to secure, consistent with Bert's physicalism, the causal connection between the content of a psychological state and its attendant behavior. Neat. Just what you would expect of Bert.

But too simple. Bert reads widely and could not help but come across Putnam's [1975] tales from Twin Earth which tell that Doppelgängers in different environments might share mental representations with identical physical and syntactic properties that nevertheless differ in content. So, Bert cast about for a division between two types of psychological content: wide (or unrestricted) and narrow (or restricted) [Putnam, 1975; Fodor, 1987, 1982, 1980; Baker 1987a, pp. 17—20]. Psychological states sharing all the same semantic properties — including truth conditions and, thus, reference — have the same wide content. But psychological states differing in selected semantic properties might yet possess the same narrow content. The distinction is not obscure. Two agents in competition for a single prize might each say, and so perhaps believe, "I will win the prize." What each believes refers to a different person and has a different truth value. So, the wide contents of their thoughts must differ. Still, there is certainly an important sense in which what they think is the same, which is why, in the relevant contexts, they act much the same. There must, therefore, be a psychologically important sense in which the contents of their thoughts coincide, and that is the sense captured by assigning them thoughts with the same narrow, restricted, psychological content. Difference in reference or truth conditions is, then, consistent with identity of narrow content.

Narrow, not wide, content contributes to explanations of behavior. Only narrow content, then, need supervene on the syntactic structure of mental representations. So, as a matter of methodology, Bert's physicalistic psychology is properly solipsistic or individualistic: narrow content is what counts in the explanation of an agent's behavior, and it varies independently of elements in the physical environment. If the narrow content of a psychological state is germane to the purposes of psychological explanation, then psychology might well ignore an agent's environment in individuating the agent's psychological states. Bert calls these wrinkles 'pleats' and says that everything is still neat.

Baker enthusiastically endorses the idea that beliefs play a fundamental role in psychological explanation, and she is equally willing to accept the idea that the proper individuation of beliefs must respect their narrow content [Baker 1987a, p. 19]. What, however, worries Baker [1987a, p. 41; pp. 88—90] is Bert's fusion of physicalism and solipsism. Bert supposes that beliefs with all the same syntactic properties are bound to be beliefs with the same narrow content. Bert is led to this view because he, like Baker, recognizes that neither the truth conditions nor the referent, if any, of a belief contributes to determining its content for the purposes of psychological explanation. Bert, still with Baker, holds that specifying a belief by its narrow content is a way of referring to a belief that respects its causal role in producing behavior. So, beliefs with the same narrow content will need to have much the same causal powers. But causal powers are physical powers. Thus, there must be, Bert reasons, some way in which the narrow content of a belief coincides with some physical feature of the belief. But the physical properties of a belief are limited to those physical properties enjoyed by whatever physical structure internal to the agent actually encodes, realizes or instantiates the belief. This is where Baker's enthusiasm for Bert's program evaporates. For Baker worries that accounts which individuate beliefs soley by reference to states internal to agents are bound crucially to fail to respect differences in the narrow contents of beliefs. She is convinced that the narrow — psychologically operant — content of an agent's belief might well depend in part on the agent's linguistic community. Baker notes [1987a, p. 37] that once Bert insists that narrow content supervenes on the syntactic, it follows that agents that are physical Doppelgängers must be syntactic duplicates as

well. If they are, then — by Bert's lights — they must enjoy psychological states with the same narrow content. Yet, Baker argues, this last need not be so. She offers a class of examples purporting to demonstrate that differently situated Doppelgängers might be in psychological states with distinct narrow contents [Baker, 1987a, p. 98]. This, Baker urges, indicates that any solipsistic or individualistic reduction of narrow content to the physical is impossible. Hence, Bert's elevated blood pressure. He cannot see any other way to accommodate the fact that psychological states are so causally connected to the environment and behavior as to respect the content of those states. Not even his paper clips can cheer Bert.

II

Bert, not Baker, is right. In order to see why, we need first to consider the example, with its associated arguments, that Baker offers to refute Bert's solipsistic reduction of narrow content to syntax. Baker [1987a, pp. 30 ff.] describes a pair of Doppelgängers whose mental representations are syntactically the same, syntactic identity being alternately gauged by physical or functional identity. Despite this coincidence, the example suggests that the agents are so socially situated that they must be accorded psychological states with divergent narrow content.

To wit: N and E, subjects in a psychological experiment, speak different, but syntactically congruent, natural languages. E speaks English; N's language is morphologically and phonetically like English. Additionally, the translation of N's language into English is largely, though not totally, homophonic. We observers distinguish between three types of assault: simple, provoked and aggravated. One striking difference between the overt languages of N and E is that E's language, but not N's, has a term for aggravated assault. Both languages recognize provoked assault, but the label for provoked assault in N's language happens to be 'aggravated assault'. When N and E both label an assault with 'aggravated', the sentence N uses attributes a provoked assault, whereas E's expression ascribes an aggravated assault. We present N and E with the same stimulus, a videotape of a violent episode, and instruct them to classify it as an assault. Both react by labeling the stimulus 'aggravated assault'. Since both subjects are competent speak-

ers of their languages, and we generally know how to translate what N says into English, Baker maintains that we must assign to N and E beliefs whose narrow contents differ. One (narrowly) pertains to provoked assault, the other to aggravated assault.

Since N and E are Doppelgängers, all their relevant behavior is nonintentionally and causally indistinguishable [Baker, 1987a, p. 31]. From this Baker reasons that, "since the bodily movements of the experimental subjects are of the same physical type and the proximate causes of the movements are of the same type, and their languages have the odd relation [that they do] . . . , it is possible that, considered physically, the mental state tokens that constituted their episodes of . . . reasoning are tokens of the same physical type" [1987a, p. 32]. She concludes that that narrow content need not supervene on the physical, at least not in the way of type-type identity [1987a, p. 32]. And, should Baker be right, neither could identity of narrow content be fixed by syntactic identity of mental representations when syntactic identity is measured by the representations functional/causal role. For the supposed physical identity of the Doppelgängers mental representations suffices, even if it is unnecessary, for the representations' syntactic identity. So, Baker's example apparently shows that N and E deploy mental representations that are syntactically identical while differing in their narrow content [Baker, 1987a, pp. 34—36].

Such is Baker's case. The argument apparently refutes Bert's solipsistic reduction of narrow psychological content to its physical or syntactical realization. For Baker aims to show that mental representations can indeed be *ambiguous* with respect to narrow content which, if correct, entails that the reduction of the narrow content of a mental representation to its syntactic form cannot be correct.

III

Two theses fuel Baker's argument. Bert had better attend to both and reject the second. The first is that, regardless of whether syntactic identity of mental representation is physically or functionally determined, N and E deploy representations of the same syntactic type. The second is that N's and E's designated beliefs differ in narrow content.

The first first. We continue provisionally to suppose with Baker that

N and E do have beliefs with different narrow content. Now, if this is correct and, as Baker agrees [1987a, p. 25], narrow content figures in explanations of behavior, the difference in the narrow contents of the beliefs *could* — even if in actual fact it does not — eventuate in behavioral differences. It is precisely this that grounds the fact that if different agents believe differently, then — other things being equal — they will, in some contexts, act differently. And, equally, if we know something about the histories of different agents, we may be entitled to attribute different beliefs to them if we catch them behaving differently in the same situation. Now, behavioral differences, whether actual or possible, are realizable in bodily movements, which movements can be described nonintentionally. Two kinds of cases want consideration: the behavioral differences might emerge in movements with some (relevant) different nonintentional descriptions; alternatively, the behavioral differences could be instantiated in movements conforming to all the same nonintentional descriptions.

Perhaps, then, difference in nonintentional descriptions separate some of the behavioral movements of N and E that are contingent upon their designated mental representations. In this event, it is utterly impossible that the mental representations encoding the different narrow contents be, as Baker requires, functionally the same. And if they are not functionally the same, they could not be physically the same. For these representations cause the movements, and yet they could not cause the differences in the movements they do if they were functionally or physically indistinguishable representations. Accordingly, if Baker's example is to tell against either the simple type-type identity thesis or Bert's functionalism, the possible behavioral differences insured by the divergence in the narrow contents of N's and E's beliefs (that Baker insists on) must be behavioral differences that are realized in movements conforming to all the same nonintentional descriptions. The unaided eye will not be able to detect any differences in N's and E's behavior.

Now the second contentious thesis. Presumably, N and E potentially differ in their behavior caused by the physically/functionally indistinguishable mental representations that realize their supposedly different narrow contents. But now these behavioral differences are rooted in bodily movements that are nonintentionally indiscernable. Of course,

the example itself illustrates just such behavior. N and E produce responses to the stimulus that are nonintentionally the same. Shown the videotape of the violent episode, both N and E might utter in the same way, "That is aggravated assault." But N is supposed, thereby, to be *asserting* that a case of provoked assault has occurred, whereas E differently *asserts* that aggravated assault has occurred [Baker, 1987a, pp. 88 f., 98—99].

Various points clamor for attention. It is no accident that all of the actual and potential behavior of both N and E caused by their respective tokens of their designated beliefs are movements of the same nonintentional types. They must be identical movements when nonintentionally described since, by Baker's lights, they are assumed to have precisely the same nonintentional causal histories. When we focus on this fact, we ought seriously to question the second of the two theses central to Baker's argument and ask whether we are entitled, finally, to attribute to N and E different belief contents. What, if anything, sanctions attributions of different belief contents to N and E is the supposition that they behave differently. In particular, we are told to take N to assert one thing and E another. But, evidently, what we take them respectively to *assert* is a function of what belief contents we have antecedently and, for the purposes of Baker's argument, *circularly* assigned to them! That is, we cannot appeal to differences in assertions to establish intentional differences in behavior instantiated in physically indistinguishable movements unless we first suppose that the behavior results from beliefs with different contents. And it is exactly that which awaits proof.

Is this backwards? Do we rather first determine what people assert and, on that basis, settle what they believe? No. What we do is note that a speaker's *utterance*, being as it is an utterance in a language with a known semantics, (probably) means such and so. This we take as evidence that the speaker believes something or other. It is only after we have assigned a belief to the speaker that we proceed to suppose that the speaker *asserts* what we take her to assert. To see that this is how the process runs it helps to attend to cases in which the meaning of the utterance and the assertion diverge because of what belief we take an utterance to indicate. Suppose that I utter 'I certainly am the murderer.' You know that this utterance means that I certainly am the

murderer. You also know that I am given to sarcasm. So, you attribute to me the belief that I am certainly *not* the murderer and, consequently, take me to assert that I certainly am *not* the murderer by way of my having sarcastically uttered something that means that I certainly am the murderer. Determining, then, what one asserts is contingent upon first fixing what one believes.

Defending the idea that the relevant beliefs of N and E disagree in narrow content, Baker takes it to be a straightforward fact that N and E are fluent, if not native, speakers of their distinct languages. This is to secure that their verbal expressions determine the contents of their assertions [1987a, pp. 31—32; 98—100]. Yet in order to see that this is incorrect note first off that whereas verbal expressions (within) a spoken language can be ambiguous, thoughts never are.[1] An ambiguous utterance cannot, then, alone determine the contents of the belief from which it arises. Hence, verbal expressions do not in general determine the contents of the mental states they may indicate.

Beyond this, suppose — consistent with the example — that either or both N and E are also perfectly fluent in the language of the other. Recall, also, that their public languages are both syntactically indistinguishable and, for the most part, homophonically translatable. Assume additionally that the audiences of N and E are known by N and E equally to include speakers of both languages. Or perhaps N and E know that their audiences are perfectly bilingual with respect to the languages in question. Under these additional assumptions, it is not at all evident from their utterances alone exactly what either N or E asserts when uttering 'That is aggravated assault' since it is now not plain which languages N and E might be using. What we need to know in order to establish what each asserts is the narrow content of what she believes, and that — as remarked — is what is to be determined. Thus, we need not endorse Baker's claim, predicated on her treatment of their assertions, that the narrow contents of N and E differ while the mental representations realizing their contents are physically the same [Grice, 1957, 1968, 1969; Schiffer, 1982].

Perhaps, in support of Baker's position, one might reply that yes, if either N or E were versed in the language of the other, then there might be some question as to the narrow content of her belief. But when each is strictly monolingual, there can be no question. As Baker [1987a, p.

29] says, the thought experiment draws, on several plausible assumptions.

> First, what a sentence says depends upon what language it is in. Second, people sometimes think in words. Third, which general belief a person expresses when sincerely and comprehendingly uttering a given sentence depends upon what language the person is speaking. Fourth, just as a single physical type of ink mark may have as tokens ink marks that have different meanings in different languages, so a single physical type of audible emission may have as tokens audible emissions that have different meanings in different languages.

We have lately seen what is troubling about the first and third assumptions. The second assumption tells us that N and E believe what they do sheerly by uttering 'That is aggravated assault'; their very utterances are now supposed to be their occurrent mental representations. The prevailing hypothesis is that N and E are monolingual speakers of different languages. By the first and fourth assumptions, their syntactically indistinguishable utterances occur within different languages and have different meanings. And so, by the third assumption, the belief that N expresses differs in narrow content from the belief that E expresses when they utter as they do. And so, should the first and third assumptions be granted, it would seem, that Baker's case is set.

The problem devolves onto the second assumption. With the exception of Harman [1973], representationalists generally [Fodor, 1975], simply do not allow that thoughts, i.e. mental representations, occur as tokens of public languages. Mental representations are typically taken to be expressions in a covert language in which mental computations are couched. So, there is simply no reason to think that N's and E's beliefs are encoded in syntactically indistinguishable mental representations even if the overt expressions of their beliefs are syntactically indistinguishable tokens.

Baker and Bert must part company on Baker's assumption that verbal agents can think in overt words. Bert holds that behavior is contingent upon the causal properties of mental representations. These properties are just the syntactic properties of the representations. Remember, as a representationalist, Bert is required to explain how the intentional properties of mental representations could affect behavior as they do. His hypothesis is that the syntax of the representations both

reflects their contents and accounts for their etiological features. Surely, the syntactic properties of N's and E's utterances cannot cause these speakers to behave as they do! Their utterances are among the behavioral effects of their relevant psychological states. Consequently, the utterances cannot be the mental representations. Bert thus denies that Baker's example plausibly demonstrates that N and E issue syntactically identical mental representations with distinct narrow contents.

Anyway, attributions of beliefs are hypotheses to the best explanation of the rationality of behavior [Dennett, 1987]. As hypotheses, attributions of belief court revision as we learn more of the agents to whom the beliefs are ascribed. If we know of N and E simply that they are monolingual speakers of different but largely syntactically congruent languages, we might naively attribute to them beliefs with different narrow contents when we catch them both producing syntactically indistinguishable utterances. But as our knowledge of N and E grows, we find it necessary to recast our belief attributions. When we discover that all of their actual and potential behavioral episodes resulting from the beliefs we attribute are physically indistinguishable and that the beliefs themselves have causal histories that are physically and nonintentionally the same, we must suspect the accuracy of our original attributions of beliefs with different narrow contents.

Urging us to suppose that N and E issue mental representations that are syntactically identical, Baker asks us to imagine that we have a complete catalogue of stimuli and subsequent responses, nonintentionally described, that figure in the learning histories of both N and E. Stunningly, the histories are absolutely the same. So, we acknowledge that the psychological states that arise from and contribute to these histories are also nonintentionally the same. Given this, it seems that were N and E periodically and unknowingly to have exchanged places during their histories, their subsequent beliefs would be the same as they actually are. For we can certainly imagine that, once we understand the details of the learning histories of each, we can reproduce the stimuli involved in these histories. What beliefs would N and E have acquired had we duplicated the stimuli to which they had been exposed and then have presented each only with the stimuli that the other had actually experienced? Surely, N's and E's beliefs arising in the counterfactual situation would be just what they were in the actual case. Yet had N and E swapped learning histories, each would have acquired the

belief actually acquired by the other. Together these counterfactual findings demonstrate that, in the actual case, N and E must have beliefs with the same narrow content.

Well, why not ask N and E to settle the case? Suppose that we inform them of all that we know of their peculiar situations. Certainly, N might maintain that the narrow content of her belief is about provoked assault, and E that the narrow content of her belief is about aggravated assault. Even if N and E should themselves side with Baker's assignment of narrow contents to them, they are in the same position as are we observers of them. Their attributions of beliefs to themselves, like belief attributions generally, are hypotheses to the best explanation of their actions [Goldman, 1989]. If we both know more about psychology than do they and are also fully informed of their learning trajectories, then we are simply better positioned than are they to attribute beliefs to them. Indeed, it would be dogmatic of either N or E to refuse to revise their self-attributions in light of what they might learn from what the best psychology might say of them. And it certainly looks as if Bert's edition of the Representational Theory of the Mind has some explanatory virtue since it (perhaps alone) has a plausible explanation of how it is that narrow content can actually cause behavior [Churchland, 1981]. If the price of a plausible psychological theory is that we must hesitate in taking the testimony of N and E as final in fixing the narrow contents of their belief states, then we need not worry that the fee is excessive.

An additional reason urges the hypothesis that if the nonintentionally indistinguishable mental representations of N and E are syntactically the same, then — but in opposition to Baker — they coincide in narrow content. Suppose, wishfully, that some very bright graduate student in Artificial Intelligence, were to write a program, COPY CAT, that perfectly simulates that band of N's behavior tethered to those of her mental representations that are syntactically the same as E's. Since N and E are related as they are, COPY CAT would emulate E as well and in the same manner as it does N. Let us assume that, within the range of its simulation, COPY CAT would merit attributions of whatever beliefs are properly attributed to N since it mimics N so well. It would, for the same reason, qualify as registering the corresponding beliefs E possesses. So, since both N's and E's designated beliefs are the same in content as COPY CAT's, the symmetry and transitivity of identity

implies that the specified beliefs of N and E, certainly including those that are syntactically the same, are also the same in content.

Suppose, then, that the targeted mental representations of N and E do agree in narrow content. What, exactly, is it? Does it pertain to aggravated assault, provoked assault or neither? It would seem that, in a sense, the common narrow content could be about neither aggravated nor provoked assault. For, following Baker, these concepts are sufficiently different that beliefs pertaining to the one must differ in narrow content from beliefs pertaining to the other. Any attempt to say that the common narrow content pertained exclusively to the one rather than the other would be *ad hoc* and could, no doubt, be reversed with equal plausibility, which is no plausibility at all. Besides, COPY CAT has a belief with the same narrow content as the belief encoded in the mental representation syntactically common to both N and E. Evidently, no reason selects between aggravated and provoked assault as the content of COPY CAT's belief. So, none could point to either concept as the content of N's and E's common representation.

Additionally, as Baker rightly remarks [1987a, pp. 98—100], if we do suppose that N and E enjoy beliefs with the same narrow content, then we must recognize that when they both utter 'That is aggravated assault' in order to express their beliefs, their utterances — taken to be utterances in their respective spoken languages — could not both correctly express their beliefs. However, this flies in the face of the natural assumption that agents typically succeed in expressing their beliefs when they attempt to do so and that their relatively rare failures to say what they believe are best attributed to slips of the tongue or purely linguistic mistakes about the meanings of words. Unless Bert can indicate what is the narrow content of N's and E's beliefs and proceed to explain how either or both N and E so thoroughly fail to express their beliefs, then perhaps he is just, as Baker [1987a, p. 99] says, "whistling in the dark."

Bert is not worried that he is ignorant of the common narrow content of N's and E's beliefs. Mentalese is the distinctive language that encodes their mental representations. It well may harbor expressions not neatly translatable into any given public language. Indeed, this is assured by Baker's very own example. For E is supposed to have a belief realized by a mental representation whose narrow content

pertains to aggravated assault. N's spoken language has no term that captures the notion of aggravated assault. Hence, there is no way perfectly to translate the narrow content of E's belief into N's spoken language! Thus and with trumpets, we should anticipate that, because of the vagaries of translation generally, we may be unable exactly to state in our (current) public languages the narrow contents of Mentalese representations.

Well now, isn't this a fine kettle of fish? Bert is forced to admit that agents typically may be quite inept at precisely saying what they believe, that an agent's ascriptions of content to herself generally lack the authority that, in the standard — perhaps introspective — case, they manifestly have. Yes, there is a long philosophical tradition that, for epistemological reasons, maintains that we are well, if not perfectly, positioned introspectively to know the contents of our own beliefs, especially those that occurrently govern our normal behavior [Stich, 1978; Pollock, 1986]. But this tradition must itself be judged by the best available evidence, and the evidence is just not favorable. Nisbett and Wilson [1977] are rightly celebrated in psychological circles — and ought to be applauded by philosophers — for having experimentally demonstrated that, in the normal scheme of things, we are amazingly inept at explaining why we act as we do and that our attempts introspectively to describe the contents of those of our current psychological states that control our behavior amount to no more than relatively uninformed speculation as to what might drive us to act as we do. Experiments aimed at establishing this hypothesis now abound, but, given the reluctance of many to appreciate their significance, it may be useful to describe the flavor of the findings.

Nisbett and Wilson report that an item's order of presentation can influence a consumer's preference for that item without the consumer being aware that her knowledge of the item's order plays any role whatsoever in her selection. In one case, subjects were presented with pairs of identical items and displayed a statistically marked preference for items as a function of their order of presentation from left to right. The subjects' knowledge of the positions of the items controlled their selection of the items. And yet, when asked to explain their choices, the subjects typically adverted to qualitative differences they claimed to have discerned among the items and vehemently denied, when subse-

quently informed of the findings, that information regarding the position of the items had in any way influenced them.

Lackner and Garrett [1972] are acknowledged for having demonstrated that the psychological states that bear on our linguistic behavior are often beyond our introspective reach. They equipped subjects with stereo earphones. Recordings of different sentences were played through each channel at roughly the same time. Subjects were taught to attend to one channel and quickly learned to recognize the sentences played through that channel. As one would expect, they were unable to report what sentences had been played through the unattended channel and were aware only of noise in that channel, though some indicated that they occasionally knew that some unidentifiable words had been played on the unattented channel. When an ambiguous sentence was played through the attended channel concurrent with a disambiguating sentence in the unattended channel, subjects displayed a significant tendency to favor interpretations of the ambiguous sentence consistent with the meaning of the disambiguating sentence. The subjects were unable to report that they had heard a sentence in the unattended channel, much less what that sentence was or expressed. Yet, they must have assigned an interpretation to the sentence; otherwise, it would not have influenced their interpretation of the ambiguous sentence. So, here again we have an instance in which a contentful psychological state plays a crucial role in determining an agent's behavior despite the fact that the agent is unable to express, indeed is unaware of, the content of that state.

Such experiments trade in ordinary behavior and cognitive processing and suggest that ascriptions of beliefs to one's self are much like ascriptions of beliefs to others. They are hypotheses to the best explanation of behavior and, as such, are subject to all the epistemic tributations typically attendant upon hypotheses compounded by problems of translation. What is important to Bert in all this is that cognitive psychology generally welcomes the idea that belief-like states, states with narrow content, do stand center stage in psychological explanation [Fodor, 1987; Stich, 1983] although we may be inept at expressing the content of such states, even in the first person. This is just Bert's position, and he likes the company he is forced to keep.

Well, maybe this is too strong. Perhaps our pronouncements regarding which of our beliefs happen to *cause* our actions are unreliable. We

might not know which of our beliefs cause us to act as we do, but we can be taken at our word when we try to say what is on our minds.[2] It is, I suppose, fundamental to the notion of belief that beliefs are mental states that play a certain kind of causal role in behavior [Stich, 1982]. Now, if we are unreliable at introspectively specifying the contents of those mental states that fulfill this causal role, we would thereby seem to be unreliable when we try introspectively to specify the contents of our *beliefs*. If, however, we are unreliable informants regarding which of our mental states induce our actions, then it is dubious that the mental states that we do identify by introspecting their contents are among our beliefs.

It is, perhaps, instructive in the present context to note the degree to which we fail at expressing in our spoken languages the contents of our mental representations. Perception, on almost all accounts, involves the mental representation of one's environment. But who is confident that one's reports about the perceived colors of a changing sunset, the sounds of a warming orchestra or the aroma in the kitchen before Sunday dinner accurately characterize the contents of one's perceptual states? As we all know, the contents of our mental representations often outstrip our linguistic resources with the result that we can construe ascriptions of narrow content to ourselves and others as only approximations, not transparent specifications, of our representations. Recognizing, then, that unreflective ascriptions of content to the psychological states of ourselves and others are subject to the fallibility of ascriptions generally, we cannot fail to appreciate that Bert's difficulty — his present inability confidently and precisely to fix the narrow content of N's and E's beliefs — is not a function of his representationalism or his infatuation with its attendant individualistic solipsism. Rather, it is endemic to any psychological theory, physically individualistic or not, that endorses the idea, one from which neither Baker nor Bert shrinks, that psychological states do indeed have narrow content.

Go ahead, Bert, put a quarter in the blood pressure machine. You'll like the new numbers.

NOTES

* I am indebted to both Lynne Rudder Baker and George Graham for helpful, critical correspondence on an earlier draft of this paper. John Pollock and Keith Lehrer

generously commented on that same draft. Conversations with Alvin Goldman, Joseph Tolliver and Juanita Simpson have been helpful on the matter of narrow content. None of these philosophers is liable for those of my mistakes that have resisted instruction.

[1] That is, a thought — as occurrent in a person — induces the person to act as she might in accordance with the content of the thought. Were thoughts ambiguous — still as occurrent tokens — the actions they induce could not generally accord with the divergent contents of the thought.

[2] John Pollock, Alvin Goldman, Lynne Rudder Baker and George Graham have all lobbied on behalf of this point.

REFERENCES

Baker, Lynne Rudder. 1987a. Saving Belief (Princeton: Princeton University Press).

———. 1987b. "Content by Courtesy," Journal of Philosophy, LXXXIV: 197—213.

———. 1986. "Just What Do We Have in Mind, in Peter A. French, Theodore E. Uehling, Jr. and Howard K. Wettstein (eds.). Midwest Studies in Philosophy: Studies in the Philosophy of Mind, 10 (Minneapolis: University of Minnesota Press), 25—48.

———. 1985. "A Farewell to Functionalism," Philosophical Studies, 48: 1—13.

Burge, Tyler. 1986. "Individualism and Psychology," Philosophical Review, 95: 3—46.

———. 1982. "Other Bodies," in A. Woodfield (ed.), Thought and Object (Oxford: Clarendon Press), 97—120.

———. 1982. "Two Thought Experiments Reviewed," Notre Dame Journal of Formal Logic, 23: 284—242.

———. 1979. "Individualism and the Mental," in P. A. French, T. E. Uehling, and H. K. Wettstein (eds.), Midwest Studies in Philosophy: Studies in Metaphysics, 4: 73—122.

Churchland, Paul M. 1981. "Eliminative Materialism and the Propositional Attitudes," Journal of Philosophy, LXXVIII: 67—90.

Dennett, Daniel. 1987. The Intentional Stance (Cambridge, Mass.: MIT/Bradford Press).

Fodor, J. A. 1987. Psychosemantics (Cambridge, Mass.: MIT/Bradford Press).

———. 1982. "Cognitive Science and the Twin-Earth Problem," Notre Dame Journal of Formal Logic, 23: 98—118.

———. 1980. "Methodological Solipsism Considered as a Research Strategy in Cognitive Psychology," Behavioral and Brain Sciences, 3: 63—109 (including peer review).

———. 1975. The Language of Thought (New York: Thomas Crowell).

Goldman, Alvin. [typescript]. "Interpretation Psychologized."

Grice, H. P. "Utter's Meaning and Intentions," Philosophical Review, 78: 147—177.

———. 1968. "Utters's Meaning, Sentence Meaning and Word Meaning," Foundations of Language, 4: 225—242.

———. 1957. "Meaning," Philosophical Review, 66: 377—388.

Harman, Gilbert. 1973. Thought (Princeton: Princeton University Press).

Lackner, J. R. and Garrett, M. F. 1972. "Resolving Ambiguity: Effects of Biasing Context in the Unattended Ear," Cognition, 1: 359—372.

Maloney, J. Christopher. 1989. The Mundane Matter of The Mental Language (Cambridge: Cambridge University Press).

Nisbett, Richard E. and Wilson, Timothy DeCamp. 1977. "Telling More Than We Can Know: Verbal Reports on Mental Processes," Psychological Review, 84: 231—259.

Pollock, John. 1986. Contemporary Theories of Knowledge (Totowa: Rowman and Littlefield).
Putnam, Hilary. 1975. "The Meaning of 'Meaning'," in his Mind, Language and Reality, 2 (Cambridge: Cambridge University Press), 215—271.
Schiffer, Stephen. 1982. "Intention-based Semantics," Notre Dame Journal of Formal Logic, 23: 119—156.
Stich, Stephen. 1983. From Folk Psychology to Cognitive Science (Cambridge, Mass.: MIT/Bradford Press).
———. 1982. "On the Ascription of Content," in A. Woodfield (ed.). Thought and Object (Oxford: Clarendon Press), pp. 153—206.
———. 1978. "Beliefs and Subdoxastic States," Philosophy of Science, 45: 573—591.

Department of Philosophy,
University of Arizona,
Tucson, AZ 85721,
U.S.A.