PHILOSOPHICAL EXPLANATIONS

INTRODUCTION

I, too, seek an unreadable book: urgent thoughts to grapple with in agitation and excitement, revelations to be transformed by or to transform, a book incapable of being read straight through, a book, even, to bring reading to *stop*. I have not found that book, or attempted it. Still, I wrote and thought in awareness of it, in the hope this book would bask in its light. That hope would be arrogant if it weren't self-fulfilling—to face toward the light, even from a great distance, is to be warmed. (Is it sufficient, though, when light is absent, to face in the direction it would emanate from?)

Familiar questions impel this essay: Does life have meaning? Are there objective ethical truths? Do we have free will? What is the nature of our identity as selves? Must our knowledge and understanding stay within fixed limits? These questions moved me, and others, to enter the study of philosophy. I care what their answers are. While such other philosophical intricacies as whether sets or numbers exist can be fun for a time, they do not make us tremble.

Our various questions stem from one: how are we valuable and precious? Consider the issue of free will, for example. Often, philosophers treat this as a question about punishment and responsibility: how can we punish someone for an action, or hold him responsible, if it was causally determined, eventually by factors going back to

before his birth, hence outside his control?* My concern with free will, however, is not rooted in a desire to punish people or hold them responsible, or even to be held responsible myself. Without free will we seem diminished, merely the playthings of external causes. Our value seems undercut. The various questions arise from, are shaped and made vivid by, a concern with our value, significance, importance, stature, and preciousness. If our lives cannot have meaning, if we are no more than puppets of causes, if our attempts at knowledge are foredoomed to failure, if we have no worth that the actions of others ought to respect, then we are devoid of value.

My concern is not only intense but directed. I want (to be able) to conclude that we *are* worthwhile and precious. But this bias does not mean I refuse to follow philosophical reason where it leads. Fortunately, two factors help me avoid conclusions of valuelessness. No philosophical argument forces us to accept its (unpleasant) conclusion; instead, we always can pursue the philosophical task of uncovering the argument's defects. This is the way Berkeley's arguments, and skeptical arguments generally, have been treated. Or we can try to find a route (believing it exists although it has not yet been found) to something almost as good as what the argument seemed to eliminate; such has been the response to Hume on induction. The second factor is an optional stop rule. I do not stop the philosophical reasoning until it leads me where I want to go; then I stop.

This description oversimplifies. In the course of the philosophical quest, the destination gets modified somewhat. Nevertheless, a value criterion is at work. The goal is getting to a place worth being, even though the investigation may change and deepen the idea of worth.

Are other philosophers more dispassionate about the important questions? A philosopher's concerns are exhibited within his work on a topic as well as in selecting that topic. When a philosopher sees that premisses he accepts logically imply a conclusion he has rejected until now, he faces a choice: he may accept this conclusion, or reject one of the previously accepted premisses, or even postpone the decision about which to do. His choice will depend upon which is greater, the degree of his commitment to the various premisses or

* I do not know of a way to write that is truly neutral about pronoun gender yet does not constantly distract attention—at least the contemporary reader's —from the sentence's central content. I am still looking for a satisfactory solution.

the degree of his commitment to denying the conclusion. It is implausible that these are independent of how strongly he wants certain things to be true. The various means of control over conclusions explain why so few philosophers publish ones that (continue to) upset them. I do not recall any philosopher reporting in distress that on some fundamental question he is forced to conclude that the truth is awful, worse even than the third best way he would want it. (Did not even Schopenhauer come to relish his conclusions?)

We may wonder whether a philosophy with a foregone conclusion can have any value at all. That a philosophy aims at a conclusion, though, does not guarantee it will reach it. We learn something of value in discovering that our goal can be reached. Since also the general aim will not determine the precise character of the conclusion reached, and since the very process of reaching may be worthwhile itself, we need not devalue a philosophical inquiry that is teleologically directed.

No philosophical inquiry can restrict itself to the central questions; in pursuing these, we are led to others as well. Common themes unite our consideration of the diverse questions but, rather than begin with these as first principles, I prefer to let linkages emerge. Philosophers often seek to deduce their total view from a few basic principles, showing how all follows from their intuitively based axioms. The rest of the philosophy then strikes readers as depending upon these principles. One brick is piled upon another to produce a tall philosophical tower, one brick wide. When the bottom brick crumbles or is removed, all topples, burying even those insights that were independent of the starting point.

Instead of the tottering tower, I suggest that our model be the Parthenon. First we emplace our separate philosophical insights, column by column; afterwards, we unite and unify them under an overarching roof of general principles or themes. When the philosophical structure crumbles somewhat, as we should expect on inductive grounds, something of interest and beauty remains standing. Still preserved are some insights, the separate columns, some balanced relations, and the wistful look of a grander unity eroded by misfortunes or natural processes. We need not go so far as to hope that the philosophical ruin, like some others, will be even more beautiful than the original. Yet, unlike the philosophical tower, this structure will remain as more than a heap of stones.

Coercive Philosophy

It was not arbitrary of philosophers to start with apparently necessary first principles, given their desire to prove their views. If the tower-like structure is abandoned, we must forsake its purpose as well. But can philosophy have an aim other than proof? Philosophical training molds arguers: it trains people to produce arguments and (this is part of the arguing) to criticize and evaluate them. A philosopher's seriousness is judged by the quality of his arguments.[1]

Children think an argument involves raised voices, anger, negative emotion. To argue with someone is to attempt to push him around verbally. But a philosophical argument isn't like that—is it?

The terminology of philosophical art is coercive: arguments are *powerful* and best when they are *knockdown*, arguments *force* you to a conclusion, if you believe the premisses you *have to* or *must* believe the conclusion, some arguments do not carry much *punch*, and so forth. A philosophical argument is an attempt to get someone to believe something, whether he wants to believe it or not. A successful philosophical argument, a strong argument, *forces* someone to a belief.

Though philosophy is carried on as a coercive activity, the penalty philosophers wield is, after all, rather weak. If the other person is willing to bear the label of "irrational" or "having the worse arguments", he can skip away happily maintaining his previous belief. He will be trailed, of course, by the philosopher furiously hurling philosophical imprecations: "What do you mean, you're willing to be irrational? You shouldn't be irrational because . . ." And although the philosopher is embarrassed by his inability to complete this sentence in a noncircular fashion—he can only produce reasons for accepting reasons—still, he is unwilling to let his adversary go.

Wouldn't it be better if philosophical arguments left the person no possible answer at all, reducing him to impotent silence? Even then, he might sit there silently, smiling, Buddhalike. Perhaps philosophers need arguments so powerful they set up reverberations in the brain: if the person refuses to accept the conclusion, he *dies*. How's that for a powerful argument? Yet, as with other physical threats ("your money or your life"), he can choose defiance. A "perfect" philosophical argument would leave no choice.

What useful purpose do philosophical arguments serve? Do we, trained in finding flaws in history's great arguers, really believe argu-

ments a promising route to the truth? Does either the likelihood of arriving at a true view (as opposed to a consistent and coherent one) or a view's closeness to the truth vary directly with the strength of the philosophical arguments? Philosophical arguments can serve to elaborate a view, to delineate its content. Considering objections, hypothetical situations, and so on, does help to sharpen a view. But need all this be done in an attempt to prove, or in arguing?

Why are philosophers intent on forcing others to believe things? Is that a nice way to behave toward someone? I think we cannot improve people that way—the means frustrate the end. Just as dependence is not eliminated by treating a person dependently, and someone cannot be forced to be free, a person is not most improved by being forced to believe something against his will, whether he wants to or not.[2] The valuable person cannot be fashioned by committing philosophy upon him.

So don't look here for a knockdown argument that there is something wrong with knockdown arguments, for the knockdown argument to end all knockdown arguing. It will not do to argue you into the conclusion, even in order to reduce the total amount of presentation of argument. Nor may I hint that I possess the knockdown argument yet will not present it.

Mightn't there be a legitimate use of argument, in self-defense against argumentative bludgeoning by others? Could one wield arguments to attack the other person's position, but only after he has attacked your own—intellectual karate in response to his initiating argument? Alternatively, arguments might be used solely to disarm an attacker. Deftly, the force of the assault could be diverted or even turned against the attacker—intellectual judo or aikido. Perhaps others could thus be defended from the onslaught of third parties, though it would be difficult to bring our argumentative defense to their attention without thereby subjecting them to coercion from *our* arguments. For one's own protection it should not be necessary to argue at all, merely to note publicly what bludgeoning the others are attempting—intellectual satyagraha, to use Gandhi's term for nonviolent resistance.*

* Is this the explanation of why philosophy department audiences try especially to refute or poke holes in lectures? The lecturer is trying to ram an opinion into their minds, so quite appropriately the audience resists, because even if it is something they want to believe anyway, they don't want to allow themselves to be *forced* to believe it.

Since a successful proof or argument must utilize premisses the listener accepts, in what way does it proceed against the person's will? In arguing for a conclusion, are you not just showing him what he already implicitly accepts? While autonomy is not violated by argument as it would be by (involuntary) brain surgery that instilled belief, still, to argue for a conclusion is not merely to point out the consequences of statements which it turns out the person believes; it is to search and cast about for suitable premisses, for statements the person does accept that will lead him to the desired conclusion. Recognizing the deductive connection, the person may either accept the conclusion or reject one of the statements he previously accepted, now that he sees where it leads. So the arguer will seek premisses the person will not abandon. His goal is not simply to point out connections among statements but to compel belief in a particular one.*

I find I usually read works of philosophy with all defenses up, with a view to finding out where the author has gone wrong. Occasionally, after a short amount of reading, I find myself switched to a different mode; I become open to what the author has to teach. No doubt the voice of the author plays a role, perhaps also his not being coercive. An additional factor affects my stance. Sometimes a writer will begin with a thought similar to one I have had and been pleased with, except that his is more profound or subtle. Or after reading the first few sentences I may have thoughts or objections which the author then will go on to state or meet more acutely. Here, clearly, is someone from whom I can learn.

I am willing to accept thoughts I read when I have had similar ones myself; I am more willing to accept my own thoughts than those I read. Yet, having a new thought is not an action intended in advance; we don't set ourselves to have that specific new thought. If our own thoughts just "come to us" unbidden, why should we be less receptive to ones that come through reading? Perhaps because we spontaneously have only those thoughts to which we are already receptive. This may lead me to miss out on learning from those who can teach me most, those who think in a way completely different from mine. Unfortunately, however, my trust in them cannot grow in

* "If a person is wondering whether or not to believe p, can't we offer him reasons to believe it as *help*?" Yes, if your help is neutral. But do you also offer reasons for not believing p? Do you pursue with further reasons for p if the first fail to convince?

the way described, so I continue to read them from an adversary stance.

Philosophy without arguments, in one mode, would guide someone to a view. The first chapter presents thoughts the reader has had (or is ready to have), only more deeply. Reading this chapter stimulates new thoughts which, pleased with, he tentatively adopts as his own. The second chapter deepens and extends these very thoughts; the reader willingly accepts them in this form. They are almost exactly what he was thinking already; he does not have to be argued into them. This second chapter also stimulates further new thoughts, which please the reader, and he tentatively adopts these thoughts as his own; in the third chapter, he finds these thoughts deepened and extended, and so on.

At no point is the person forced to accept anything. He moves along gently, exploring his own and the author's thoughts. He explores together with the author, moving only where he is ready to; then he stops. Perhaps, at a later time mulling it over or in a second reading, he will move further.

With this manner of writing, an author might circle back more than once to the same topic. Not everything can be said at once or twice; a reader may not be ready yet to think it all himself.[3] Within the structure of each chapter, the thought might go further out as it goes along, reaching finally ideas so speculative that even the author is not willing (yet?) to assert them, barely willing even to entertain them.

Such a book could not convince everybody of what it says—it wouldn't try. (Should it then be judged by goals not its own?) I have said such a book would guide without forcing, but won't it be manipulating its readers? Not every way a teacher can help someone to see something himself, more deeply, counts as manipulation, especially when the activity is acknowledged mutually.

This is not my prelude to an announcement, even though I would like to present a philosophical view in this way, author and reader traveling together, each continually spurting in front of the other. Not only do I lack the art to do this, I do not yet have a philosophical view that flows so deeply and naturally. Perhaps a philosophy should be tested by whether it can be presented so.

Even this least dominating mode of writing maintains a hierarchy wherein the reader is to be attentive to the writer's thoughts but not

conversely. (Though the writer must attend to the unknown readers' possible responses, this does not alter the asymmetry.) Do some contemporary French critics propose to deconstruct texts in order to destroy that hierarchy, held to be undesirable, to destroy the very distinction between writer and reader? It would be ironic if a writer's desire not so to dominate his readers were to lead him to withhold thoughts. (How can thoughts or worries about the writer–reader or speaker–listener asymmetry be conveyed without invoking that very asymmetry?) It is only the writer's thoughts, though, that can be rejected and dismissed—is it because asymmetrical attentiveness is matched by asymmetrical vulnerability that anything (voluntarily) gets written and read?

I place no extreme obligation of attentiveness on my readers; I hope instead for those who read as I do, seeking what they can learn from, make use of, transform for their own purposes. Much as they wanted to be understood accurately, the philosophers of the past would have preferred this response, I think, to having their views meticulously and sympathetically stated in all parts and relations. The respect *they* paid their predecessors was philosophy, not scholarship. Rather than our listening to them, wouldn't they prefer we spoke to them? (We have to listen closely enough, though, to speak to *them*.)

Philosophical Explanations

There is a second mode of philosophy, not directed to arguments and proofs: it seeks explanations. Various philosophical things need to be explained; a philosophical theory is introduced to explain them, to render them coherent and better understood.

Many philosophical problems are ones of understanding how something is or can be possible. How is it possible for us to have free will, supposing that all actions are causally determined? Randomness, also, seems no more congenial; so, how is free will (even) possible? How is it possible that we know anything, given the facts the skeptic enumerates, for example, that it is logically possible we are dreaming or floating in a tank with our brain being stimulated to give us exactly our current experiences and even all our past ones? How is it possible that motion occurs, given Zeno's arguments? How is it

possible for something to be the same thing from one time to another, through change? How is it possible for subjective experiences to fit into an objective physical world? How can there be stable meanings (Plato asked), given that everything in the world is changing? How is it possible for us to have synthetic necessary knowledge? (This last question, Kant's, shows, if none did earlier, that the question's presupposition that the item *is* possible may be controversial or even false, in which case the question would be withdrawn.) The theological problem of evil also takes this form: how is evil possible, supposing the existence of an omnipotent omniscient good God? One central question of twentieth century philosophy has been: how is language possible? And let us not omit from our list: how is philosophy possible?

The form of these questions is: how is one thing possible, given (or supposing) certain other things? Some statements $r_1, \ldots, r_n$ are assumed or accepted or taken for granted, and there is a tension between these statements and another statement p; they appear to exclude p's holding true. Let us term the r_i *apparent excluders* (of p). Since the statement p also is accepted, we face the question of how p is possible, given its apparent excluders.* Note that the question is not: given p, how are the apparent excluders possible? Tension and incompatibility are symmetrical relations among statements, yet typically philosophical problems focus on the possibility only of some statements on one side of the relation. It is an interesting issue, what determines in which direction the question is salient.

The strongest mode of exclusion would be logical incompatibility: the apparent excluders, in conjunction, logically (appear to) imply that p is false; they imply the negation of p, which we may write as not-p.

Philosophical arguments show or draw upon (apparent) incompatibilities or other tensions between statements; these can have an im-

* A physical analogue of a philosophical problem is the familiar puzzle of the pencil or rod with a nondetachable closed loop of string attached to one end of it, a loop too small to reach also around the other end. Placed in a buttonhole in the appropriate way, attached to it, it quickly comes to seem impossible to remove, *provably* impossible—the loop is shorter than the length of the rod. Yet you don't doubt that it can be removed, since another person has done it again and again; the question is how this is possible. (The solution also provides an illuminating analogue to one way of solving a philosophical problem, of explaining how something is possible.)

portant role in setting-up philosophical problems without being designed to force belief.

Given the (apparent) incompatibility between the apparent excluders and *p*, there are two ways to continue to maintain (the possibility of) *p*. First, one of the apparent excluders can be denied, or there can be a denial of their conjunction all together. To save the possibility of *p*, it is not necessary to prove these denials, only to show we need not accept one of the apparent excluders or their conjunction. Second, each of the apparent excluders can continue to be maintained, while their apparent incompatibility with *p* is removed, either by close scrutiny showing the reasoning from them to not-*p* to be defective, or by embedding them in a wider context or theory that specifies how *p* holds in the face of these apparent excluders.

To rebut an argument for not-*p* from specific apparent excluders removes a reason for thinking *p* cannot hold, and so counts as a kind of explanation of how *p* can be possible. This task is unending, for as knowledge advances, or seems to, new apparent excluders come to the fore, and hence new questions arise about the possibility of *p*. "If we know that whenever a new apparent excluder comes along, we will try to show that *p* remains standing, wouldn't it be more economical simply to prove *p* once and for all?" This proposal misconstrues the need. A proof of *p* will give us the conviction that *p* is true, but it need not give us understanding of how *p can* be true (given the apparent excluder). Even when the argument from an apparent excluder does not lead us to deny *p* or to doubt its truth, it still may leave us puzzled as to how *p* can be true. Typically, the arguments of the epistemological skeptic do not lead us to conclude we don't have knowledge; but they do leave us wondering how we can know what we do. A proof that *p* is true, however, need not show how *p* is compatible with the apparent excluders, or show which apparent excluder is false—it need not mention them at all. So the task of showing how *p* is possible cannot be done once and for all by a proof that *p*. What a proof can do—show us that *p* is true—is not what we need, for we already believe this. Why isn't it enough to know that *p* is true, why do we also need to understand how it can be true? To see how *p* can be true (given these apparent excluders) is to see how things fit together. This philosophical understanding, finding harmony in apparent tension and incompatibility, is, I think, intrinsically valuable. Yet I would not try to bludgeon anyone into needing or wanting it.

The task of explaining how p is possible is not exhausted by the rearguard action of meeting arguments from its apparent excluders. There remains the question of what facts or principles might give rise to p. Here the philosopher searches for deeper explanatory principles, preferably with some independent plausibility, not excluded by current knowledge. To show that these principles, if true, would explain p involves deducing p from them—at least so holds the deductive-nomological view according to which each explanation deduces the fact to be explained from general laws and initial conditions.[4] Yet still, this is no attempt to prove p; and the explanatory hypotheses used in the explanation need not be known to be true, or be believed on grounds independent of p itself.

To produce this possible explanation of p is, by seeing one way p is given rise to, to see how p can be true. "How is it possible that p? This way: such and such facts are possible and they constitute an explanatory route to p." The more true-like these explanatory hypotheses, the more we understand how p can be true. The (possible) explanation of p from them is put forward tentatively, subject to withdrawal in the face of difficulties or alternative, better explanations, perhaps using deeper principles that also would explain other things.[5]

Status of the Hypotheses

Which hypotheses may be introduced in a philosophical explanation? May the hypothesis be known to be false? I believe increased understanding can be produced even by an explanation known to be false; seeing what in principle could give rise to a phenomenon illuminates some of its aspects by the way it latches onto these. Richard Feynman notes[6] that the inverse square character of gravitation would follow if particles flew equally in all directions but were partially blocked by bodies, so fewer reached the earth sunward than from the other side (and fewer hit the sun earthward than from the other side), so that each body was propelled toward the other by the hitting particles. Yet this potential explanation is known to be false; no effect is observed of the moving earth's being slowed down by the larger number of particles hitting it from the front. Still, even though false, this explanation produces some illumination and increased understanding.

I am tempted to say that explanation locates something in actuality, showing its actual connections with other actual things, while *understanding* locates it in a network of possibility, showing the connections it would have to other nonactual things or processes. (Explanation increases understanding too, since the actual connections it exhibits are also possible.) Recall how illuminating it can be to place something, something actual even, in a typology or a two-by-two matrix, how salient is the insight gained through locating it in that network of alternative possibilities. No less understanding is provided by an explanatory hypothesis which might coexist with the phenomenon and generate it. When this book explores hypotheses depicting eccentric possibilities, as it sometimes does, even a reader who is convinced the hypothesis fails, who will not take the possibility seriously, even a reader who does not enjoy (as I do) the playful exploration of possibilities for its own sake, may see benefit in the increased understanding gained.

However much it may increase understanding, a hypothesis known to be false will not explain how something *is* possible. (If we can find no other explanation, though, we may reconsider whether we know the hypothesis is false.) In addition to not being known to be false, must the hypothesis also be plausible? I do not think so, even leaving aside how fragile and parochial are judgments of plausibility. The question of how p is possible may cut so deeply that the only answers which suffice are implausible, at least as one judges before investigating how p is possible. Moreover, insisting the hypothesis be plausible antecedently would guarantee that philosophy could not lead us to radically new and surprising truths or insights. One should not rule out even rejecting the p begun with, if intensive effort fails to show how it is possible while showing how something similar is.[7]

Does the philosopher who explains how p is possible, by putting forth potential explanations of p, differ from the scientist who puts forth and tests potential explanations of p in order to explain why p is true? I would not want to claim that philosophical explanation must be discontinuous with scientific activity. Yet typically, the philosopher's hypothesis is not testable or disconfirmable, because he puts forth only an existentially quantified statement; he says there is something or other, some process or other, that satisfies certain general structural conditions and so yields p. That there is or might be a

process of that sort shows how p is possible. To specify the particular details of a process of that sort would be to engage in empirical science: differing scientific specifications each would fit the philosopher's existential statement, which holds merely that there is some or another true specification.[8]

The epistemologist may need for his purposes only the fact that our perceptions somehow respond to presented facts so as to satisfy certain general conditions of responsiveness; to show how knowledge is possible he need only speculate on a linkage of that sort existing. To explain why our perceptions thus respond to the facts, however, is a task for the perceptual and physiological psychologists, who must specify the details of the particular mechanism whereby responsiveness is achieved, and for the evolutionary psychologist who must explain how that mechanism arose and was selected for. Still, although the philosophical and scientific activities typically differ, the philosopher's existential hypothesis may suggest detailed investigations to the scientist; conceivably the philosopher might specify the sort so completely that its existence is immediately open to empirical test.[9]

Explanation versus Proof

Philosophical argument, trying to get someone to believe something whether he wants to believe it or not, is not, I have held, a nice way to behave toward someone; also, it does not fit the original motivation for studying or entering philosophy. That motivation is puzzlement, curiosity, a desire to understand, not a desire to produce uniformity of belief. Most people do not want to become thought-police. The philosophical goal of explanation rather than proof not only is morally better, it is more in accord with one's philosophical motivation. Also it changes how one proceeds philosophically; at the macro-level (as we already have noted), it leads away from constructing the philosophical tower; at the micro-level, it alters which philosophical "moves" are legitimate at various points.

Even if (deductive) proof and (deductive) explanation have the same abstract structure, wherein p is deduced from the statements q, the pragmatics of the two activities differ. In the case of explanation, the thing (p) to be explained, to be deduced, must be known or at

least believed to be true. If you are asked to explain why you are not reading this book now, you reject the request. Since you *are* reading it now, there is no fact that you are not, to be explained. You do not set out to explain what you don't believe to be true. However, you may set out to prove what you don't (yet) believe to be true; establishing it as true will induce the belief. In order to set about proving something, you need not actually (yet) believe the conclusion p to be deduced.

In order to prove p, however, you must start from premisses q which you, or those to whom you are proving p, know or believe. Suppose I set out to prove to you that God exists by beginning with the statement "everything Maimonides believed is true"; I then cannot waive aside your objection that you do not believe that, saying I am going to proceed to prove God's existence nonetheless. A proof transmits conviction from its premisses down to its conclusion, so it must start with premisses (q) for which there already is conviction; otherwise, there will be nothing to transmit.[10] An explanation, on the other hand, may introduce explanatory hypotheses (q) which are not already believed, from which to deduce p in explanatory fashion. Success in this explanatory deduction itself may lend support and induce belief, previously absent, in the hypothesis.

The activity of explaining how something is possible is hardly new to philosophy, but we here diverge from previous views of its nature; for example, Kant required philosophy to be apodictic, using only principles that were certain and necessarily true.

Philosophical explanations, however, can be offered tentatively, the hypotheses or theories presented might be believed only tentatively or even not believed initially at all; they can be held subject to revision, or introduced for the purpose of seeing in principle how p could be explained. To propose that explanation replace proof as the goal of philosophy is not to suggest a completely new activity that rejects all previous philosophical work. That would be absurd. Many important arguments philosophers offered can be viewed as raising explanatory questions by the incompatibilities they appear to establish—hence philosophers have a continuing interest in paradoxes—while others can be recast into explanatory form so as to present possible explanations rather than purported proofs. Indeed, much of the practice of philosophers, especially recently, as opposed to their metaphilosophical talk, easily fits the explanatory mold; making the

explanatory goal explicit has the virtue of legitimating the introduction of explanatory hypotheses that are uncertain.*

The shift from the interpersonal goal of proof has the further micro-effect of altering which statements may be introduced when. Consider the philosophical problem of skepticism; this has been presented and pursued as the problem of *refuting* the skeptic,[11] of proving to him that he does know what he doubts he knows, or of proving to him that you do know what he denies you know. My concern with skepticism is different, fitting within philosophy as an explanatory activity.[12]

My purpose is not to refute the skeptic, to prove he is wrong, to convince him, to marshal arguments and reasons which must convince him (if he is rational). In being unconcerned to convince the skeptic, I may seem not to be taking what he says seriously, but in a way I am taking what he says more seriously than someone does who merely sets out to convince him. If I attempt to convince the skeptic of something p, that is a task of the foreign relations department of my belief system. I have to find something q which the skeptic accepts (it does not matter whether I also accept it) from which he will conclude that p. In a discussion of skepticism, if q is said, someone can appropriately object that the skeptic will not or need not or should not accept q, if the purpose of the discussion is to convince the skeptic or to discover what will or must convince him. In trying

* Philosophers sometimes have offered arguments or proofs, namely transcendental arguments, in which explanatory considerations were intermixed. Philosophical explanations and transcendental arguments both start from a p which is accepted, and seek explanatory hypotheses q which play a role in giving rise to p. However, a transcendental argument seeks a q which is a precondition of p, something without which p could not be true. Thereby, and that is its purpose, it argues for the truth of q; q must be true, for p is true and p could not be true unless q were. A transcendental argument begins with the question "how is p possible?", but since its goal is to *prove* something it must find a q which not only explains p's possibility (as part of a set of sufficient conditions for p) but also is a necessary condition of it. If we were concerned only to explain how p is possible, a sufficient condition would do. For the philosophical questions we shall be concerned with, it is difficult enough to conjure up sufficient conditions, without also searching for necessary ones and trying to prove they are necessary. There is a difference between explaining p via q, and proving q is the correct explanation of p. A transcendental argument attempts to prove q by proving it is part of any correct explanation of p, by proving it a precondition of p's possibility.

to convince the skeptic, what is relevant is how *his* beliefs fit together—that is why it is foreign relations.

But the attempt to explain how knowledge is or can be possible, given what the skeptic says, is a task for my belief system's bureau of internal affairs. Some of the things the skeptic says or points out (for example, that certain situations are logically possible) I accept; these are or become part of my own belief system. My problem is that I don't see (or no longer see, after the skeptic has spoken) how these things go along with yet other things in my belief system, namely, numerous beliefs that I and others know certain things. My task here is to remove the conflict, to put my own beliefs in alignment, to show how those of the things the skeptic says which I accept can be fit in with other things I accept. In this way, I take very seriously what the skeptic says, for I acknowledge that what he says creates a problem for *me* and my beliefs. In thus trying to explain to myself how knowledge is possible, what is relevant is what I accept; the explanation is no less acceptable to me because the skeptic rejects part of it. The goal of explanation makes it legitimate for the philosopher to introduce statements as hypotheses (acceptable to him) that the goal of proof would exclude as begging the question (of proving to the skeptic that he knows). Although my goal is not interpersonal, I assume that I am not idiosyncratic in how what the skeptic says presents a problem to me; others also will not see how knowledge is possible given what the skeptic says. If the explanation I offer draws only on things these other people also (can) accept, it will be acceptable to them as well, whether or not the skeptic accepts it, even though my goal is not to explain it to them.*

I might describe the situation as follows. I take the skeptic less seriously than someone does who sets out to convince him, whose concern is what the skeptic thinks or continues to think; but I take what the skeptic says more seriously than someone does who merely sets out to convince him, for I view what the skeptic says as a problem for me, for my beliefs. My problem does not disappear if the

* If, however, the skeptic's words convince you, or lead you no longer to believe you know, then you do not face the task of explaining how it is possible for you to know, and no hypothesis will gain support (in your eyes) by performing this task. Isn't it rash, though, to be convinced by the skeptic's arguments, to have more confidence that the appatent incompatibility he invokes is a real one than that you know various things?

skeptic does, or if the skeptic says he was only joking, that he really does believe he know things. If I succeed in my task, I (and others like me) learn from the skeptic, my beliefs change and are reorganized, while the skeptic need not learn anything. Whereas if the convincer succeeds, the skeptic will learn something, though the convincer himself need not. I do not take the skeptic seriously enough to want to teach him; I do take what he says seriously enough to want to learn from it. It should be emphasized that, though internal to my (nonidiosyncratic) belief system, this task is not one of self-development; the goal is to explain, to understand (in this case) how knowledge is possible.

This distinction between the foreign and domestic relations of my system of beliefs illuminates another puzzle. In discussions of ethics one sometimes thinks, "how could one convince some particular evil figure, say Stalin, Hitler, or Mao, that he is wrong; if there is no argument guaranteed to convince him, doesn't that show that ethics really is subjective, merely a matter of preference or opinion?" The puzzle is why we take this possibility of irresolvable disagreement to threaten the objectivity of ethics, while, for example, an irresolvable disagreement about whether someone in a mental institution is Napoleon or Jesus does not, we think, threaten the objectivity of such factual matters, even though there is no way we can convince the deluded person that he is wrong. I suggest the distinction lies here: the mental patient presents us with a problem of foreign relations, but what he says does not cause difficulties about how our beliefs fit together within our own belief system. Therefore, it need not concern us if some parts of our explanation of why he is not and cannot be Napoleon or Jesus are rejected by him. We have a general conception of past history, how we come to know some facts about it, why (even were a "return" possible) not everyone who claims to be Napoleon or Jesus can be right, why some people form self-important delusions, and so forth. Our general picture of the world of objective historical truths and of our connection to it leads us to think we know why the (deluded) person's claims must be wrong; we have some conception of how we can and do know he is wrong. The disagreements that worry us are those that fix upon tensions or anomalies or unclarities in our own belief system. It is because we do not see how an objective ethics is *possible* that we worry about irresolvable moral disagreements. For this reason, we do not need actually

to find such disagreement—if it doesn't exist, we will invent it. For it is the possibility of such disagreement that worries us; it worries us because we do not see clearly how there can be an objective ethics. If we did see how an objective ethics was possible, we would not attribute theoretical importance to someone who could not be brought to agree.

Thus, what is philosophically interesting, what demarcates the philosophically important disagreements from the others, is the domestic problem presented for our own beliefs. Because this is what creates the philosophical interest, it is on this explanatory issue we shall concentrate, rather than on the philosophically pointless task of attempting to convince the other person.[13]

Philosophical Pluralism

Can a philosophy begin by seeking philosophical explanation? Doesn't the desire for explanation rather than something else already presuppose philosophical views about what is intellectually desirable? And, even given the goal of explanation, won't philosophies differ in their conception of an acceptable or adequate explanation, in the conditions that they hold an explanation must meet? One view will hold that explanation must be deductive while another will not, one that the principles used in explanation must be necessary or self-evident, or that explanation must utilize general laws, or that explanations must be causal, or mechanical, or teleological, or contain a picturable model, or not refer to unobservable entities, or be poetic, or be testable by observation, or comport with revealed religious doctrines, while other views will differ. Since each philosophy will have its own ideas about explanation and when it is suitable, how can we begin philosophy by seeking explanations? Must we not already have answered many important philosophical questions or presupposed answers to them in order to fix what it is we are seeking? Similar questions to these about explanation can be asked also about the notions of 'incompatibility' and 'tension' which set up the need for certain explanations, and about the statements themselves which are held to be in tension.

These questions are not unique to the goal of explanation. To seek for proofs also has its presuppositions, and philosophies can differ in

the conditions they impose on adequate proofs. (Are they deductive, formalizable, formalizable in first-order logic, finitistic, may they proceed by reductio ad absurdum? In addition, there are the many conditions that can be imposed on the axioms or propositions from which proofs begin.) It is clear that any goal put forth for philosophy will be philosophically controversial in that some philosophies will reject that goal in favor of another, while yet other philosophies which accept it will specify it differently. How, then, is it possible for philosophy to get started?

I do not see how to satisfy the desire to start philosophy in a neutral way, making no philosophical assumptions, remaining neutral among all possible philosophical views. We cannot reasonably hope to settle on one philosophical view by showing it uniquely satisfies all of the apparently neutral desirable general conditions on a philosophy; many different philosophies will equally well satisfy those. Anyway, why think the goal of neutrality itself is neutral? Would every philosophy accept *it*? And when neutrality ends eventually, as we want, what selective factor will point to one philosophy rather than others, and what status—neutral or not—will *this* factor have? If a neutral beginning is chimerical, the alternative of starting just where we are seems parochial and dogmatic, especially if there are some theoretical places we can't get to from here.

We can build modes of change into a view, hoping parochialism is avoided when any theory can be reached, in principle, given suitable input. Nevertheless, some places will not be reached with specified inputs, starting from here; philosophical views will differ in what gets reached when. A metaphilosophy will be part of a total philosophical view rather than a separate neutral theory above the battle.

The treatment for philosophical parochialism, as for parochialism of other sorts, is to come to know alternatives. We can keep track of the different philosophical views that have been put forth and elaborated; we can pay attention to foreign traditions and their diverse viewpoints, to the special slant of these traditions on our questions, both the different ways they pose their most nearly equivalent question, and the different answers they offer.[14] There even may be ways of catapulting oneself, at least temporarily, into different philosophical perspectives. Various drugs seem to have given the experience of how the world looks and feels from one or another of the diverse Eastern perspectives to Westerners, including some people

previously unfamiliar with the conceptual framework into which they were catapulted.*

Not all philosophical cosmopolitans, keeping track of alternative positions, will end up the same, for they start differently. I stand on my position keeping track of the others, while you stand on yours. The position I occupy is modified somewhat by my broader knowledge, but I do not imagine that all the different positions so modified by knowledge of the others will converge. After all, we do not even have agreement among philosophies within the same tradition. Nor do I recommend withdrawing assent from every particular position and merely contemplating the panoply of alternatives; I do not, that is, recommend "rootless cosmopolitanism".** Assent will, though, be more tentative, perhaps more transient. (This tentative attachment to one view, while carrying along the rest, would not jibe with the goal of convincing others or proving conclusions.)

This fits my experience in studying philosophy: I confess I have found (and not only in sequence) many different philosophies alluring and appealing, cogent and impressive, tempting and wonderful. I think this says something about the subject, not about me—I am not noticeably wishy-washier than most. Treating philosophy as a black box, we might view its "output" not as a single theory, not even as one set of theories, but as a set of questions, each with its own set of associated theories as possible answers. Should we view the highest products of philosophy as the philosophical questions themselves, the theories and systems being commentary to exhibit the value of the questions? On this view, philosophy's wonders are the ones in which it begins; as important as new answers are its new questions.

I feel discomfort, though, with the aesthetic view of philosophy, the uncommitted praise of the diverse philosophical "visions". The goal is finding out the truth, after all. (Yet is that goal, or its specification, neutral among philosophies?) Recall the distinction made ear-

* The explanation of why some years ago books on Eastern thought started to appear in profusion in philosophy sections of paperback bookstores, previously almost purely Western, surely was the onset of widespread experimentation with drugs. People were turning to these books as fitting and explaining what they suddenly were experiencing unpreparedly.

** Applied to Jews throughout the world before the establishment of the state of Israel, this term confused lack of roots with lack of the land to nourish them.

lier between explaining and understanding. Embedding the world in the network of alternative philosophical theories and visions, seeing how each of these different philosophical possibilities gets a grip on the world, does produce understanding. The major philosophical theories of continuing interest are readings of possible worlds accessible from here, that is, possible readings of the actual world. We understand the world by seeing it in its matrix of possibilities, in its possibility neighborhood.

Still, this book does not aim at understanding; by and large, it aims more narrowly—at explanation, at truth. In any case, I would not think it a good strategy to aim at illuminating possibility; we are all too likely to fail in our aim. Aiming for the true explanation might well yield understanding, while striving for understanding probably will not produce even that, only inaccessible possibilities. (Such intellectual playfulness has its own value, though, and is not excluded from this book.) Moreover, my desire is to explain how knowledge is possible, how free will is possible, how there can be ethical truths, how life can have meaning. That is what I want to know. Other philosophical views are scanned and searched for help.

Yet, it seems to me I value truth less than do other philosophers I know. (Do I love truth less or love understanding more?) How could I so admire the writings of history's great philosophers if falsity were a fatal or debilitating flaw?

I see the situation as follows. There are various philosophical views, mutually incompatible, which cannot be dismissed or simply rejected. Philosophy's output is the basketful of these admissible views, all together. One delimiting strategy would be to modify and shave these views, capturing what is true in each, to make them compatible parts of one new view. While I know of no reason in principle why this cannot be done, neither has anyone yet done it satisfactorily. Perhaps, as knowing a subject (such as logic or physics) involves seeing the different ways it can be organized and viewed, the different ways around it, so too (only this time the views are incompatible so the analogy is imperfect) knowing the world involves seeing the different ways it can be viewed.

Are we reduced to relativism then, the doctrine that all views are equally good? No, some views can be rejected, and the admissible ones remaining will differ in merits and adequacy, though none is completely lacking. Even when one view is clearly best, though, we

do not keep only this first ranked view, rejecting all the others. Our total view is the basket of philosophical views, containing all the admissible views.[15] This total view notices which component is best, and perhaps will order the others. Yet the first ranked view is not completely adequate all by itself; what it omits or distorts or puts out of focus cannot be added compatibly, but must be brought out and highlighted by another incompatible view, itself (even more) inadequate alone.

The position is not relativism, for the views are ranked, but neither is just one view settled upon.[16] True, philosophers will differ in what they hold inadmissible, and in how they rank the views that remain. I am not proposing a neutral way to do this, simply describing how things look from within the view I rank first (which includes a metacomponent). Is there at least this agreement: even if there are different perspectives on the world, isn't it true on all perspectives that the world looks a particular way from a particular perspective? These relational facts are about a perspective but hold independently of a perspective. Can't we use them to figure out what the world is like, what character of world would give rise to these relational facts?[17] However, with philosophical views, there is no way to identify the point from which the world looks a certain way, other than by how it looks—the world looks as if it fits Plato's theory from the position of standing on Plato's theory. It is misleading to call the different views "perspectives", in the absence of any independent way of identifying positions.

Still, we can ask what the world and our cognitive relation to it are like so as to give rise to this situation: a basketful of admissible (partially) ranked philosophical views. We shall return to this question at the close of the book. Meanwhile, we have explored the metasituation sufficiently to get on with offering specific philosophical explanations.

I would not want, however, to leave the impression that explanation is the sole legitimate goal of philosophy. Even within the realm of truth and actuality, there are linkages, connections, and contrasts other than explanatory ones mediated by general laws; these ways of setting something in (actual) context produce understanding too. Furthermore, the broader end of understanding something in its (possible) context is valuable. Other philosophical activities also have independent legitimacy, not merely as subordinate within the

explanatory framework; the skeptic's arguing that certain things are *not* possible is an example. My pluralism extends to ways of carrying on philosophy; though I elaborate one, I keep track of and appreciate others. Even argument, despite its coerciveness and distortion of philosophical direction, has its own virtues: it is responsive to a person's rationality (if not his autonomy), and perhaps in caring what he believes it evidences some caring about the person. But though philosophical arguing is somewhat responsive to and respectful of the person, it is not responsive and respectful enough.

The view of philosophy as philosophical explanation is put forth here as a tentative hypothesis, designed to encompass much of the actual historical activity of philosophers while demarcating a legitimate and important task. Moreover, the view applies to and fits itself. In explaining how philosophy is possible, given the formidable obstacles to it as a useful mode of knowledge, the view itself is an instance of what it says philosophy should be: the explanation of how something is possible. In contrast, the view that philosophy is the theory of self-evident fundamental principles and their consequences, for example, seems neither to be self-evident nor derivable from such. No doubt, not everything in this book (and certainly not everything of value in the history of philosophy) fits the mold sketched here. Still, this view has the virtue of delineating one important direction philosophy can pursue, a direction whose nature is reasonably well understood and which does not appear beyond our capacity.

Many philosophers have dreamed of setting philosophy upon the sure path of a science. As this book's closing makes evident, this dream is not mine. Even the path of science, though, does not involve the proof of theories. It is a commonplace of the philosophy of science that evidence is incomplete, that alternative hypotheses and possibilities can be imagined, that theories are held tentatively until a better one is produced, and so on. Philosophy has aspired to more than this, as did explanatory science itself once, but surely we long ago reached the point when philosophy should cease striving for so much more while accomplishing so much less.

The explanations to follow are put forward not as the sole correct view on their topics, but as members among others of admissible classes, with the hope that they will be ranked first, or at least highly. On the view presented here, philosophical work aspires to produce a

highest ranked view, at least an illuminating one, without attempting to knock all other theories out as inadmissible. (Even this general view of the basketful with rankings need aspire only to being first ranked in its basket.)

This view of philosophy is in harmony with our earlier substitution of explanation for proof as the goal of philosophy. Yet, though proof and convincing fit best with admissible sets of only one member, they are not excluded even by the view of the full basket. Why not try, then, to convince the reader or prove to him that the explanations offered are admissible and should be highly ranked? I don't say there could *not* be occasions when it was possible and appropriate to convince someone of this, or of the fact that he was enjoying tennis, or to prove to someone that he is in love—but should we build an activity and methodology around those occasions? It is (hard) enough to offer an explanation which the reader will find illuminating. There is no need also to prove it is. (Would one prove it by presenting something the reader simply accepts as proof, or would one also have to prove that proves it?)

This book puts forward its explanations in a very tentative spirit; not only do I not ask you to believe they are correct, I do not think it important for me to believe them correct, either. Still, I do believe, and hope you will find it so, that these proposed explanations are illuminating and worth considering, that they are worth surpassing; also, that the process of seeking and elaborating explanations, being open to new possibilities, the new wonderings and wanderings, the free exploration, is itself a delight. Can any pleasure compare to that of a new idea, a new question?

There is sexual experience, of course, not dissimilar, with its own playfulness and possibilities, its focused freedom, its depth, its sharp pleasures and its gentle ones, its ecstacies. What is the mind's excitement and sensuality? What its orgasm? Whatever, it unfortunately will frighten and offend the puritans of the mind (do the two puritanisms share a common root?) even as it expands others and brings them joy.